
Working with Apache Access Logs in Sumo Logic

◆ Step 1: Basic Search with `_sourceCategory`

Query:

```
_sourceCategory=Labs/Apache/Access
```

✓ Explanation

- `_sourceCategory` is a **metadata tag** that classifies logs into categories (in this case, Labs/Apache/Access for Apache access logs).
- Running this simple search will return all raw Apache access logs in that category.
- Example log entry:
- 192.168.1.5 - - [19/Aug/2025:10:45:12 +0000] "GET /index.html HTTP/1.1" 200 2326
- This log contains important fields such as IP address, request method, URL, HTTP version, and **status code (200 in this case)**.

◆ Step 2: Parsing Status Codes

Query:

```
_sourceCategory=Labs/Apache/Access  
| parse "HTTP/1.1\" * " as status_code
```

✓ Explanation

- The parse operator is used to **extract specific values** from raw logs.
- In Apache logs, the HTTP status code appears right after "HTTP/1.1".
- The statement "HTTP/1.1\" * " tells Sumo Logic to:
 - Look for the text HTTP/1.1" in the log.
 - Extract the next value (represented by *).
 - Store it as a new field called status_code.

Example Output:

```
status_code
```

200

404

500

Now instead of just raw logs, you have a structured field (status_code) that can be counted, grouped, or filtered.

◆ Step 3: Identifying Client vs Server Errors

Query:

```
_sourceCategory=Labs/Apache/Access
```

```
| parse "HTTP/1.1\" * " as status_code
```

```
| if(status_code matches "4*", 1, 0) as client_err
```

```
| if(status_code matches "5*", 1, 0) as server_err
```

```
| sum(server_err) as server_errors_cnt, sum(client_err) as client_errors_cnt
```

✅ Explanation of Operators

1. if(condition, value_if_true, value_if_false)

- A conditional operator that assigns values based on logic.
- Example:
 - if(status_code matches "4*", 1, 0) → If the status code starts with 4 (client error), mark it as 1, else 0.
 - if(status_code matches "5*", 1, 0) → If the status code starts with 5 (server error), mark it as 1, else 0.

2. matches

- Used inside the if condition to check patterns.
- "4*" matches any status code starting with 4 (like 400, 403, 404).
- "5*" matches any status code starting with 5 (like 500, 502, 503).

3. sum(field)

- Adds up all the values in a field.

- Since `client_err` and `server_err` are either 1 or 0, summing them gives the **total number of errors**.

Final Output Example:

```
server_errors_cnt  client_errors_cnt
```

```
-----
```

```
12              85
```

Why This Matters

- **Client Errors (4xx):** Show when users make invalid requests (e.g., 404 Not Found, 403 Forbidden). High counts may indicate broken links or probing activity.
 - **Server Errors (5xx):** Show when the server fails to handle a request. High counts may indicate server issues or overload.
 - By parsing and aggregating with conditions, you transform **raw logs into actionable insights** that can highlight problems in real-time.
-